

Two Studies in Data-Constrained Deep Learning: Chest X-ray Classification and English–Urdu Neural Translation

Soban Ahmad

FAST National University of Computer and Emerging Sciences
sobanahmad2003@gmail.com

Abstract—We report two independent studies. In Q1, a custom 1.24M-parameter CNN (PneumoNet) is trained from scratch to classify chest radiographs into COVID-19, Normal, and Pneumonia, and compared against frozen and fine-tuned VGG16 and ResNet50 baselines under identical conditions. PneumoNet outperforms all four pre-trained baselines on every test metric (accuracy 0.9797, macro-F1 0.9794, macro-AUC 0.9982), including fine-tuned VGG16 at 108× the parameter count. In Q2, a vanilla RNN encoder–decoder (no attention, no gating) is trained on an English–Urdu sentence-pair corpus for neural machine translation; the originally reported result was diagnosed as under-trained rather than constraint-limited, and correcting the teacher-forcing schedule together with reversing the source sequence raises test BLEU from 2.45 to 4.64, though the corrected model remains fluent but largely unfaithful to its source. Both studies surface measurement-driven findings that contradict default assumptions in the respective pipelines: a duplicate-detection threshold imported from natural-image practice fails on chest radiographs, horizontal flipping corrupts a laterality marker, frozen backbones silently adapt through BatchNorm statistics, an L2 regularization sweep is numerically inert at float32 precision, a Q2 quality ceiling initially attributed to the no-attention architecture was actually an under-trained teacher-forcing configuration, and the Q2 corpus is a ninth of the assignment’s stated size and overwhelmingly Biblical in register.

I. INTRODUCTION

A. Q1: Chest X-ray Classification

Automated triage of chest radiographs into COVID-19, Normal, and Pneumonia classes is a natural fit for CNNs, but production pipelines default to large ImageNet-pretrained backbones without asking whether a small, task-specific network trained from scratch might do as well with far fewer parameters and no transfer-learning risk. This study builds PneumoNet, a 1,239,779-parameter CNN, and benchmarks it against frozen and fine-tuned VGG16 (134.3M params) and ResNet50 (23.5M params) under identical data, splits, and training budget. Contributions: a calibrated near-duplicate detection procedure for radiographs (Section II), a measured augmentation policy rather than an assumed one (Section IV), a BatchNorm-freezing correction that changes what “frozen

baseline” actually means (Section III), and a full grid-search-tuned comparison showing the small model wins (Section V).

B. Q2: English–Urdu Neural Machine Translation

A vanilla RNN encoder–decoder is the base case against which attention and gating are usually justified, but that justification is rarely measured directly, and rarely disentangled from ordinary training bugs. This study trains an nn.RNN-only encoder–decoder (no LSTM, GRU, attention, or Transformer component) on an English–Urdu parallel corpus and evaluates it with BLEU, length-bucketed BLEU, and an error taxonomy. Contributions: an audit of the actual corpus against the assignment’s description (Section II); a three-configuration sweep that diagnoses the originally reported quality ceiling as under-training (a static 0.5 teacher-forcing rate and forward-order source) rather than architecture, and corrects it, nearly doubling test BLEU (Section IV); and evidence that, even after correction, translation quality remains bounded not by sentence length but by the single fixed-size hidden vector’s ability to condition generation faithfully on the source (Section VI).

II. DATASET AND PREPROCESSING

A. Q1

The source is the Kaggle prashant268/chest-xray-covid19-pneumonia dataset: 6,432 images across three classes, all decodable (0 corrupt). Near-duplicate detection was recalibrated during development. An initial dHash-8/Hamming-distance- ≤ 3 rule flagged 1,570 of 6,432 images (25%) as near-duplicates; a 40-pair audit of the flagged set found zero genuinely near-identical pairs, a mean flagged-pair correlation of 0.842 against a random-pair baseline of 0.605, and 10 of 40 flagged pairs spanning two different diagnostic classes. The root cause: chest radiographs share gross anatomy and framing, so the dataset-wide median nearest-neighbour thumbnail correlation is already 0.926 (p99 0.970, p99.5 0.987) — a perceptual-hash threshold calibrated on natural images does not transfer to this domain. The rule was replaced with contrast-normalized 32×32 -thumbnail correlation ≥ 0.99 , which sits in the




Fig. 1: Heavy-policy augmentation examples. Rows 3–4 show the R laterality marker mirrored into a Cyrillic-“ya”-shaped artifact under horizontal flip.

genuine-duplicate tail and removes 34 exact plus 15 near duplicates, 49 of 6,432 images (0.76%). One of the removed near-duplicate pairs crosses diagnostic classes: PNEUMONIA(7).jpg (dropped) and COVID19(277).jpg (kept) correlate at 0.9913 — the same radiograph filed under two different diagnoses in the raw data, measured label noise that deduplication can silence but not resolve, since which of the two original labels was correct is unknowable from the image alone. The kept manifest (6,383 images: COVID19 560, Normal 1,579, Pneumonia 4,244) is split 80/10/10, stratified by class, giving train/val/test = 5,106/638/639 images, a 7.6:1 Pneumonia-to-COVID19 imbalance carried into a weighted loss (Section III). Augmentation policy was measured, not assumed: Fig. 1 shows that horizontal flipping mirrors the radiopaque R laterality marker into a mirrored, Cyrillic-“ya”-shaped artifact — not just anatomically questionable for a chest film, but a visible, class-irrelevant shortcut cue an augmented model could exploit.

B. Q2

The source is the Kaggle muhammadnoman76/translation-dataset English–Urdu spreadsheet. The assignment describes a corpus of approximately 24,000 general parallel sentence pairs; the file actually delivered contains **9,103** pairs in a single sheet. Word-level auditing of the raw English side shows the corpus is overwhelmingly Biblical text

in King James register: “unto” appears in 22.7% of sentences, “god” 13.1%, “shall” 11.0%, “jesus” 10.4%, “thou”/“thee”/“thy” in 6.8%/4.3%/3.4%, over 6,561 unique English word types across 187,640 tokens — extremely repetitive, with a small modern-conversational tail late in the file. Preprocessing NFC-normalizes and lowercases English, folds Arabic-form Urdu characters to Urdu forms (deliberately excluding alef-with-madda, a distinct Urdu letter, from that folding), strips diacritics, and separates punctuation into its own tokens. Cleaning drops 1 empty and 10 exact-duplicate pairs (9,092 kept), and a length filter (95th-percentile per-side token cap, max length ratio 2.5) removes a further 627 long and 25 skewed-ratio pairs, leaving 8,440 pairs split 80/10/10 (6,752/844/844, unstratified — no labels to stratify on). Word-level vocabularies (min. frequency 2) give 3,656 source and 3,961 target types, with validation OOV rates of 3.18% (English) and 3.50% (Urdu).

III. METHODOLOGY

A. Q1

PneumoNet (Table I, 1,239,779 params) is four convolutional blocks doubling channels $32 \rightarrow 64 \rightarrow 128 \rightarrow 256$, each block two Conv-BatchNorm-ReLU [3] pairs followed by max-pooling, then global average pooling and a dropout-regularized two-layer classifier head.

TABLE I: PneumoNet architecture: layer types, output shapes and parameter counts.

Layer	Output Shape	Parameters
PneumoNet	[1, 3]	1,239,779
Sequential: 1-1	[1, 256, 14, 14]	1,173,216
Sequential: 2-1	[1, 32, 112, 112]	10,208
Conv2d: 3-1	[1, 32, 224, 224]	864
BatchNorm2d: 3-2	[1, 32, 224, 224]	64
ReLU: 3-3	[1, 32, 224, 224]	—
Conv2d: 3-4	[1, 32, 224, 224]	9,216
BatchNorm2d: 3-5	[1, 32, 224, 224]	64
ReLU: 3-6	[1, 32, 224, 224]	—
MaxPool2d: 3-7	[1, 32, 112, 112]	—
Sequential: 2-2	[1, 64, 56, 56]	55,552
Conv2d: 3-8	[1, 64, 112, 112]	18,432
BatchNorm2d: 3-9	[1, 64, 112, 112]	128
ReLU: 3-10	[1, 64, 112, 112]	—
Conv2d: 3-11	[1, 64, 112, 112]	36,864
BatchNorm2d: 3-12	[1, 64, 112, 112]	128
ReLU: 3-13	[1, 64, 112, 112]	—
MaxPool2d: 3-14	[1, 64, 56, 56]	—
Sequential: 2-3	[1, 128, 28, 28]	221,696
Conv2d: 3-15	[1, 128, 56, 56]	73,728
BatchNorm2d: 3-16	[1, 128, 56, 56]	256
ReLU: 3-17	[1, 128, 56, 56]	—
Conv2d: 3-18	[1, 128, 56, 56]	147,456
BatchNorm2d: 3-19	[1, 128, 56, 56]	256
ReLU: 3-20	[1, 128, 56, 56]	—
MaxPool2d: 3-21	[1, 128, 28, 28]	—
Sequential: 2-4	[1, 256, 14, 14]	885,760
Conv2d: 3-22	[1, 256, 28, 28]	294,912
BatchNorm2d: 3-23	[1, 256, 28, 28]	512
ReLU: 3-24	[1, 256, 28, 28]	—
Conv2d: 3-25	[1, 256, 28, 28]	589,824
BatchNorm2d: 3-26	[1, 256, 28, 28]	512
ReLU: 3-27	[1, 256, 28, 28]	—
MaxPool2d: 3-28	[1, 256, 14, 14]	—
AdaptiveAvgPool2d: 1-2	[1, 256, 1, 1]	—
Sequential: 1-3	[1, 3]	66,563
Flatten: 2-5	[1, 256]	—
Dropout: 2-6	[1, 256]	—
Linear: 2-7	[1, 256]	65,792
ReLU: 2-8	[1, 256]	—
Dropout: 2-9	[1, 256]	—
Linear: 2-10	[1, 3]	771
Total		1,239,779

The four baselines are ImageNet-pretrained VGG16 [1] and ResNet50 [2], each run frozen (only a new head trained) and fine-tuned (VGG16: last convolutional block unfrozen; ResNet50: `layer4` unfrozen), so ten total configurations reduce to the five reported after selecting one setting per architecture pairing shown in Table V. A correction was required for the frozen baselines: `requires_grad=False` blocks gradient updates to BatchNorm affine parameters but does not stop BatchNorm from updating its `running_mean/running_var` buffers in `train()` mode, since that update is unconditional on `requires_grad`. On `resnet50_frozen`, one measured forward/backward/optimizer step moved `bn1.running_mean` by 0.1565 before the fix and exactly 0.0 after forcing frozen BatchNorm modules to stay in `eval()` mode across every `train()` call. Without this fix, "frozen" baselines would have silently adapted their BatchNorm statistics to the chest X-ray domain every epoch, inflating their scores and understating the frozen-vs-fine-tuned gap Section V reports.

B. Q2

The encoder-decoder [4] (Table II, 4,770,425 params) uses `nn.RNN` with tanh nonlinearity exclusively — no LSTM, GRU, attention module, or Transformer component anywhere in the model, a hard constraint of the study rather than a design preference, so that Section VI can characterize what a single fixed-size hidden vector can and cannot carry.

TABLE II: Q2 vanilla RNN encoder-decoder: layer parameter shapes (src_vocab=3656, tgt_vocab=3961).

Layer	Parameter Shapes	Parameters
encoder.embedding	(3656, 256)	935,936
encoder.rnn	(512, 256), (512, 512), (512,), (512,)	394,240
decoder.embedding	(3961, 256)	1,014,016
decoder.rnn	(512, 256), (512, 512), (512,), (512,)	394,240
decoder.fc_out	(3961, 512), (3961,)	2,031,993
Total		4,770,425

The encoder packs its embedded input with `pack_padded_sequence (enforce_sorted=False)` before the recurrence so PAD steps never contaminate the final hidden state handed to the decoder. Decoder inputs and targets are built under an explicit shifting contract: for target tokens $[w_1, w_2]$, `dec_in` = `[SOS,  $w_1, w_2$ , PAD...]` and `dec_target` = `[ $w_1, w_2$ , EOS, PAD...]`, with a masked cross-entropy loss (Section IV) ignoring PAD positions on the target side.

IV. EXPERIMENTAL SETUP

A. Q1

All models train with AdamW [7], a class-weighted loss (inverse label frequency) for the 7.6:1 imbalance, and early stopping on validation macro-F1. Hyperparameters were selected by staged coordinate descent over 30 proxy configurations at 128×128 resolution (stages A: augmentation $\times$ normalization, B: batch size $\times$ lr, C1-C3:

dropout, L2, L1, D: epochs $\times$ patience), not an exhaustive Cartesian product — a full grid over these axes would run into days of MPS wall time, which the project's compute budget did not allow. Table III reports the searched range and selected value for every axis.

TABLE III: Grid search: hyperparameter, range searched, and optimal value.

hyperparameter	range	optimal
augmentation	heavy, light, none	none
normalization	dataset, imagenet, zero_one	imagenet
batch_size	16, 32, 64	16
lr	0.0001, 0.0003, 0.001	0.0001
dropout	0.2, 0.3, 0.5	0.3
l2_lambda	0, 1e-05, 0.0001	0
l1_lambda	0, 1e-05	0
epochs	30, 60	60
patience	5, 10	10

Two results are reported honestly rather than smoothed over. First, augmentation: stage A shows augmentation monotonically *hurting* validation macro-F1 across all three normalization schemes without exception (none > light > heavy; best none+imagenet 0.9745, worst heavy+dataset 0.9508), consistent with Section II's finding that this dataset is clean, consistently framed, and harmed rather than helped by geometric invariance training. Second, L2 weight decay: the searched range $\{0, 10^{-5}, 10^{-4}\}$ returned bit-identical validation macro-F1 (0.9756799123668505) for all three values in stage C2. This is not a discreteness artifact of the metric; an isolated 200-step AdamW experiment at the same learning rate reproduces it directly — $\lambda \in \{0, 10^{-5}, 10^{-4}\}$ give identical loss and weights, while $\lambda = 10^{-2}$ differs. AdamW's decoupled update $p \leftarrow p - \eta\lambda p$ at $\eta = 10^{-4}$, $\lambda = 10^{-4}$ is a relative per-step change of 10^{-8} , below float32 epsilon ($\approx 1.19 \times 10^{-7}$), so it rounds to exactly no change: the searched range was below representable precision, a trap for any small-weight-decay/small-learning-rate float32 experiment. The searched range and optimal value ($\lambda = 0$) are reported as-is in Table III. The search's own top selection for epochs/patience was 60/10 (validation macro-F1 0.9761, 49 epochs run), narrowly ahead of 30/5 (0.9757, 17 epochs) — a gain of 0.0004 for roughly $3\times$ the compute. Under a stated compute budget, the five full-size runs reported in Section V used 30/5, not the search's 60/10 selection; both numbers are reported here so the near-tie between configurations $3\times$ apart in cost stands as a tuning observation in its own right, not a discrepancy to hide. All five full runs, including VGG16 fine-tuning, completed at the winning batch size of 16 with no MPS out-of-memory event, so the pre-registered batch-size fallback was not needed. Because Metal (MPS) kernels are not bitwise-deterministic, splits, vocabularies, and every selected hyperparameter reproduce exactly across runs, while metric values reproduce to roughly $\pm 0.3\%$; the seed is 42 throughout, set once in a shared `set_seed` utility.

B. Q2

Training minimizes masked cross-entropy (PAD positions excluded) with Adam and gradient-norm clipping at 1.0. Clipping is genuinely active early in training: 16% of batches were clipped in epoch 1 (mean pre-clip norm 0.815, max 1.702), falling below 1% of batches by epoch 2 (mean pre-clip norm 0.693, max 1.068) — evidence the vanilla RNN has a real, if modest, gradient-explosion tendency that clipping is actively controlling, not a ceremonial safeguard.

The originally reported model (test BLEU 2.45) was diagnosed as *under-trained*, not constraint-limited: it could not reproduce training pairs it had already seen 32 times, and `train_loss` (4.297) exceeded `val_loss` (4.244) — a positive gap characteristic of underfitting, not the capacity ceiling a hidden-vector bottleneck would produce. Two configuration causes, not architectural ones, were identified and fixed. First, `teacher_forcing` was fixed at 0.5 from epoch 1, so the decoder consumed its own untrained predictions half the time and never received a clean training signal. Second, the source sequence was fed in forward order; reversing it [4] shortens the dependency between the first source and first target words, the standard remedy for an attention-free RNN encoder-decoder. Table IV reports all three measured configurations. Correcting teacher forcing alone *lowered* test BLEU (2.45→1.56) even though it improved validation perplexity (68.08→57.76); only combined with source reversal did BLEU nearly double, to 4.64, with perplexity falling further to 53.37 — perplexity and BLEU did not move together across this sweep, worth noting honestly rather than smoothing into a single trend line. `reverse_source` changes token order only: it adds no LSTM, GRU, attention, or Transformer component, so the vanilla-RNN constraint stated in Section III still holds for the reported (final) configuration. Under that corrected configuration, training ran 18 epochs to an early-stopping halt (configured budget 80/patience 8), 4,770,425 parameters, 210.3s total.

TABLE IV: Q2 tuning sweep: teacher forcing and source order (val. perplexity, test BLEU).

Configuration	Val. Perplexity	Test BLEU
TF=0.5, forward source (orig.)	68.08	2.45
TF=1.0, forward source	57.76	1.56
TF=1.0, reversed source (final)	53.37	4.64

V. RESULTS

A. Q1

Table V is the full cross-model test-set comparison. PnuemoNet leads on every metric: accuracy 0.9797, macro-F1 0.9794, macro-AUC 0.9982, against fine-tuned VGG16 (0.9765/0.9759/0.9980, 134.3M params) and fine-tuned ResNet50 (0.9734/0.9753/0.9949, 23.5M params). Frozen baselines trail badly — VGG16-frozen macro-F1 0.9160, ResNet50-frozen 0.9407 — confirming ImageNet features transfer weakly onto grayscale radiographs without adaptation.




Fig. 2: PneumoNet test-set performance.




Fig. 3: Test accuracy / macro-F1 vs. training-set fraction. Saturates at 75%.

Fig. 2 shows the normalized confusion matrix, ROC curves, and training curves for PnuemoNet. The data-volume study (Fig. 3) trains PnuemoNet on 25/50/75/100% of the training data: test macro-F1 is 0.9517 (25%, $n = 1,277$), 0.9670 (50%, $n = 2,554$), 0.9794 (75%, $n = 3,829$), and 0.9776 (100%, $n = 5,106$) — performance saturates at 75%; the final quarter of training data adds nothing, and the 100% result is marginally lower, within run-to-run noise, not a monotonic climb.

B. Q2

The final configuration (teacher forcing 1.0, reversed source) reaches best validation loss 3.9772 (perplexity 53.37) and test BLEU [6] 4.64 over 844 held-out pairs — an A/B/C tuning result rather than a single run (Table IV, Section IV). Fig. 4 shows training/validation loss for the final configuration, across its 18 epochs, falling well below the untrained baseline $\ln(3,961) \approx 8.28$.

Length-bucketed test BLEU (Table VII, fresh samples in `results/q2_samples.csv`) is close to flat across source lengths, unlike the earlier configurations. Qualitatively, the model produces fluent, register-appropriate Urdu at both short and long source lengths, but that fluency does not imply faithfulness. Table VI shows the pattern: the model reproduces Biblical register and anchors on individual content words, yet does not compose a faithful translation. In the first example the opening clause is reproduced exactly before the predicate diverges; in the

TABLE V: Cross-model comparison on the held-out Q1 test split.

model	params (total)	params (trainable)	epoch time (s)	test acc.	test macro-F1	test macro-AUC
pneumonet	1239779	1239779	50.1662	0.9797	0.9794	0.9982
vgg16_frozen	134272835	12291	61.4905	0.9249	0.9160	0.9836
vgg16_finetune	134272835	7091715	72.3212	0.9765	0.9759	0.9980
resnet50_frozen	23514179	6147	34.6194	0.9405	0.9407	0.9925
resnet50_finetune	23514179	14970883	48.4463	0.9734	0.9753	0.9949




Fig. 4: Q2 training/validation loss, final configuration, 18 epochs.

second a single content word survives while the rest of the proposition is lost. Both patterns recur throughout `results/q2_samples.csv`.

VI. ERROR ANALYSIS AND DISCUSSION

A. Q1

Only 13 of 639 test images are misclassified. By true class: COVID19 0, Normal 5, Pneumonia 8 — zero errors on the rarest class (56 test images), evidence the inverse-frequency-weighted loss is doing its job rather than being swamped by the majority Pneumonia class. Confusion is concentrated in the Normal/Pneumonia boundary (Pneumonia→Normal 7, Normal→Pneumonia 5, Pneumonia→COVID19 1). The model is also appropriately uncertain when wrong: mean confidence is 0.853 on errors versus 0.989 on correct predictions, not confidently wrong. Fig. 5 shows the 12 highest-confidence errors (of 13 total). These results are only honest because of the BatchNorm fix in Section III: pre-fix, the frozen baselines would have partly adapted to the domain via drifting running statistics, inflating their scores and narrowing the frozen-vs-fine-tuned gap this section reports as real. On regularization, dropout (0.3) did the substantive work of controlling overfitting in the winning configuration, while L2 was numerically inert over its searched range (Section IV) and L1 measurably *hurt* validation macro-F1 (0.9757 at $\lambda = 0$ vs. 0.9665 at $\lambda = 10^{-5}$) — for this model and this amount of data, explicit weight penalties added little beyond what dropout and early stopping already provided.

B. Q2

The error taxonomy (`results/q2_error_taxonomy.json`) reproduce training pairs it has already seen repeatedly over 844 test pairs, final configuration, shows UNK rate




Fig. 5: Twelve highest-confidence Q1 misclassifications.

12.88%, repetition rate 3.79%, empty-output rate 0%, and mean length ratio (hypothesis/reference) 1.0544 — the model now over-produces slightly relative to the reference and still rarely degenerates into empty output. Both the UNK rate and the repetition rate rose sharply from the original, under-trained run (2.46%→12.88% and 0.36%→3.79% respectively) even as test BLEU nearly doubled (Section IV): correcting under-training was not a uniform improvement across metrics, and a configuration that wins on corpus-level n-gram overlap can measurably lose on token-level degeneracy.

Length-bucketed BLEU is close to flat: 1–5 tokens 4.343, 6–10 tokens 3.721, 11–15 tokens 3.315, 16+ tokens 4.125 (Table VII). This flattening is not a surprise but a confirmation: source reversal shortens the dependency between early source tokens and early target tokens (Section IV), so it should help most exactly where the fixed-size hidden vector previously had the furthest information to carry, i.e. the longer buckets — which is what the earlier length-degradation pattern reflected and reversal has now removed.

A flat length curve does not mean the fixed-size-hidden-vector limitation is gone; the evidence for it is now different, and stronger. First, faithfulness: even the best configuration produces fluent, domain-appropriate Urdu that is frequently unfaithful to the specific source content rather than merely awkward (Section V examples) — consistent register, but wrong or absent propositional content. Second, memorization: the model still cannot reliably reproduce training pairs it has already seen repeatedly during training (32 times in the original diagnostic run, 18

TABLE VI: Two test-set translations from the final configuration (teacher forcing 1.0, reversed source). Glosses are literal back-translations of the hypothesis.

Source	Reference	Hypothesis (gloss below)
blessed are the meek for they shall inherit the earth	ہونگے وارث کے زمین وہ کیونکہ ہیں حلیم جو وہ ہیں مبارک	ہے نزدیک کے خدا جو وہ ہیں مبارک “blessed are those who are near to God” — opening clause exact, predicate wrong
and he touched her hand and the fever left her	گئی اتر سے پر اس تپ اور چھو ہاتھ کا اس نے اس	ہے دبا جھنوائے کس کو اس کہ کہا آکر پاس کے اس وہ اور “and he came to him and said, who touched him?” — retains <i>touched</i> , loses the fever and the healing

TABLE VII: Q2 final-configuration test BLEU by source-length bucket.

Source length (tokens)	BLEU
1–5	4.343
6–10	3.721
11–15	3.315
16+	4.125

times in the final configuration), which is not a length or vocabulary-coverage problem since these exact pairs have already been encountered many times over; it is a sign that a single fixed-size hidden vector is not conditioning generation strongly enough on the specific input, which is exactly the bottleneck attention was introduced to remove [5]. The Q2 corpus’s narrow, repetitive Biblical domain (Section II) is precisely why this is visible: it lets the model fall back on register-specific n-gram regularities that sound right and score reasonably even when the content is wrong, so fluency and raw BLEU understate the faithfulness gap, and neither BLEU 4.64 nor the flat length curve should be read as evidence the bottleneck has been resolved.

VII. CONCLUSION AND FUTURE WORK

A. Q1

A 1.24M-parameter CNN trained from scratch, under a measured preprocessing and augmentation policy and a corrected BatchNorm-freezing implementation, outperforms fine-tuned VGG16 (108× the parameters) and fine-tuned ResNet50 (19× the parameters) on every held-out metric, with only 13 total misclassifications on the test set. Limitations: the dataset is single-source and only 6,383 images after deduplication, chest radiographs across classes are intrinsically visually correlated (Section II), and no external or multi-institution validation was performed, so generalization beyond this dataset’s imaging protocol is unverified. Future work: multi-site external validation, class-activation visualization (e.g. Grad-CAM) to check the model attends to lung fields rather than markers or positioning artifacts, and re-running the augmentation study with anatomically valid transforms (small rotations, intensity jitter) rather than horizontal flips.

B. Q2

Once an under-trained teacher-forcing schedule and a forward-only source order were diagnosed and corrected (test BLEU 2.45→4.64; Section IV), the vanilla RNN

encoder–decoder still shows the limitation this study set out to measure, just via different evidence than originally reported: even its best configuration produces fluent, domain-appropriate Urdu that is frequently unfaithful to the source and still cannot reliably reproduce training pairs it has seen repeatedly, indicating the single fixed-size hidden vector it must pass from encoder to decoder under-conditions generation on the specific input (Section VI) rather than merely truncating long sentences. Limitations: the corpus is 9,103 pairs, not the assignment’s stated ~24,000, and is domain-narrow and repetitive (Section II), so these results should not be read as general-purpose English–Urdu translation quality; the corrected configuration also raised the UNK and repetition rates even as BLEU improved, so the fix is a net gain, not a uniform one. The named remedy, evidenced rather than assumed, is architectural: adding an attention mechanism [5] so the decoder can access all encoder states instead of one bottleneck vector, or replacing `nn.RNN` with a gated unit (LSTM/GRU) to improve long-range gradient flow, would directly target the faithfulness gap this study measured. Future work also includes sourcing a larger, general-domain parallel corpus so results are comparable to the published NMT literature.

REFERENCES

- [1] K. Simonyan and A. Zisserman, “Very Deep Convolutional Networks for Large-Scale Image Recognition,” *ICLR*, 2015.
- [2] K. He, X. Zhang, S. Ren, and J. Sun, “Deep Residual Learning for Image Recognition,” *CVPR*, 2016.
- [3] S. Ioffe and C. Szegedy, “Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift,” *ICML*, 2015.
- [4] I. Sutskever, O. Vinyals, and Q. V. Le, “Sequence to Sequence Learning with Neural Networks,” *NeurIPS*, 2014.
- [5] D. Bahdanau, K. Cho, and Y. Bengio, “Neural Machine Translation by Jointly Learning to Align and Translate,” *ICLR*, 2015.
- [6] K. Papineni, S. Roukos, T. Ward, and W.-J. Zhu, “BLEU: a Method for Automatic Evaluation of Machine Translation,” *ACL*, 2002.
- [7] I. Loshchilov and F. Hutter, “Decoupled Weight Decay Regularization,” *ICLR*, 2019.